

```

from __future__ import absolute_import, division, print_function, unicode_literals

import pathlib


import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns

try:
    # %tensorflow_version only exists in Colab.
    %tensorflow_version 2.x
except Exception:
    pass
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers

print(tf.__version__)

```

 Colab only includes TensorFlow 2.x; %tensorflow\_version has no effect.  
2.11.0

```

import pandas as pd
from sklearn.datasets import load_boston
boston = load_boston()
print(boston.data.shape) #get (number of rows, number of columns or 'features')
print(boston.DESCR) #get a description of the dataset

```

```

(506, 13)
.. _boston_dataset:

```

Boston house prices dataset

-----

**\*\*Data Set Characteristics:\*\***

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B  $1000(Bk - 0.63)^2$  where Bk is the proportion of black people by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference on

/usr/local/lib/python3.8/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function load\_boston is deprecated; `load\_

The Boston housing prices dataset has an ethical problem. You can refer to

The documentation of this function can be found here

```
# Next, we load the data into a 'dataframe' object for easier manipulation, and also print the first few rows in order to examine it
data = pd.DataFrame(boston.data, columns=boston.feature_names)
data.head() #notice that the target variable (MEDV) is not included
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
#For some reason, the loaded data does not include the target variable (MEDV), we add it here
data['MEDV'] = pd.Series(data=boston.target, index=data.index)
data.describe() #get some basic stats on the dataset
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	4.038095	2.299738	312.628463	18.461316	393.310082	6.339977
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.149739	1.612947	81.651513	2.161924	57.856234	9.270373
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	2.120000	1.000000	222.000000	15.200000	392.830000	4.030000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	4.030000	2.000000	242.000000	17.800000	396.900000	4.980000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	4.038095	2.299738	312.628463	18.461316	393.310082	6.339977
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	4.090000	2.300000	322.000000	18.700000	396.900000	7.900000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	4.090000	3.000000	711.000000	21.000000	396.900000	9.140000

```
data.tail() #check out the end of the data (last 5 rows)
```


	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88

```
data.isna().sum()
```

CRIM	0
ZN	0
INDUS	0
CHAS	0
NOX	0
RM	0
AGE	0
DIS	0
RAD	0
TAX	0
PTRATIO	0
B	0
LSTAT	0
MEDV	0
dtype:	int64

```
train_dataset = data.sample(frac=0.7,random_state=0)
test_dataset = data.drop(train_dataset.index)

sns.pairplot(train_dataset[["MEDV", "CRIM","AGE","DIS","TAX"]], diag_kind="kde")
```



```
train_stats = train_dataset.describe()
train_stats.pop("MEDV")
train_stats = train_stats.transpose()
train_stats
```

	count	mean	std	min	25%	50%	75%	max
CRIM	354.0	3.767375	9.418497	0.00906	0.082757	0.274475	3.077295	88.9614
ZN	354.0	11.079096	23.070178	0.00000	0.000000	0.000000	12.500000	95.0154
INDUS	354.0	11.185254	6.646944	0.74000	5.860000	9.795000	18.100000	27.7400
CHAS	354.0	0.070621	0.256554	0.00000	0.000000	0.000000	0.000000	1.0000
NOX	354.0	0.554098	0.115748	0.38500	0.453000	0.538000	0.624000	0.8710
RM	354.0	6.265791	0.699380	3.56100	5.878250	6.175000	6.605500	8.7790
AGE	354.0	68.057627	27.953167	6.00000	45.100000	76.500000	93.750000	100.00
DIS	354.0	3.844439	2.187514	1.12960	2.073700	3.207450	5.214600	12.1270
RAD	354.0	9.440678	8.569207	1.00000	4.000000	5.000000	20.000000	24.0000
TAX	354.0	407.500000	162.296676	187.00000	287.000000	337.000000	666.000000	711.00
PTRATIO	354.0	18.461299	2.149735	12.60000	17.325000	18.850000	20.200000	22.0000
B	354.0	352.720650	95.764288	2.60000	373.852500	390.945000	396.225000	396.9000
LSTAT	354.0	12.614011	7.020224	1.73000	7.347500	11.185000	16.635000	37.9310

```
train_labels = train_dataset.pop('MEDV')
test_labels = test_dataset.pop('MEDV')

def norm(x):
    return (x - train_stats['mean']) / train_stats['std']
```

```
normed_train_data = norm(train_dataset)
normed_test_data = norm(test_dataset)

def build_model():
    model = keras.Sequential([
        layers.Dense(64, activation='relu', input_shape=[len(train_dataset.keys())]),
        layers.Dense(64, activation='relu'),
        layers.Dense(1)
    ])

    optimizer = tf.keras.optimizers.RMSprop(0.001)

    model.compile(loss='mse',
                  optimizer=optimizer,
                  metrics=['mae', 'mse'])
    return model
```

```
model = build_model();
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 64)	896
dense_1 (Dense)	(None, 64)	4160
dense_2 (Dense)	(None, 1)	65
=====		
Total params: 5,121		
Trainable params: 5,121		
Non-trainable params: 0		

```
example_batch = normed_train_data[:10]
example_result = model.predict(example_batch)
example_result
```

```
1/1 [=====] - 0s 347ms/step
array([[ 0.09018825],
       [-0.0524627 ],
       [ 0.30674988],
       [ 0.10789406],
       [ 0.21827169],
       [ 0.0014355 ],
       [ 0.17999958],
       [-0.08920994],
       [ 0.33306706],
       [ 0.6581632 ]], dtype=float32)
```

```
# Display training progress by printing a single dot for each completed epoch
```

```
class PrintDot(keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs):
        if epoch % 100 == 0: print('')
        print('.', end='')
```

```
EPOCHS = 1000
```

```
history = model.fit(
    normed_train_data, train_labels,
    epochs=EPOCHS, validation_split = 0.2, verbose=0,
    callbacks=[PrintDot()])
```

.....

.....

.....

.....

.....

.....

.....

.....

.....

.....

```
hist = pd.DataFrame(history.history)
hist['epoch'] = history.epoch
hist.tail()
```

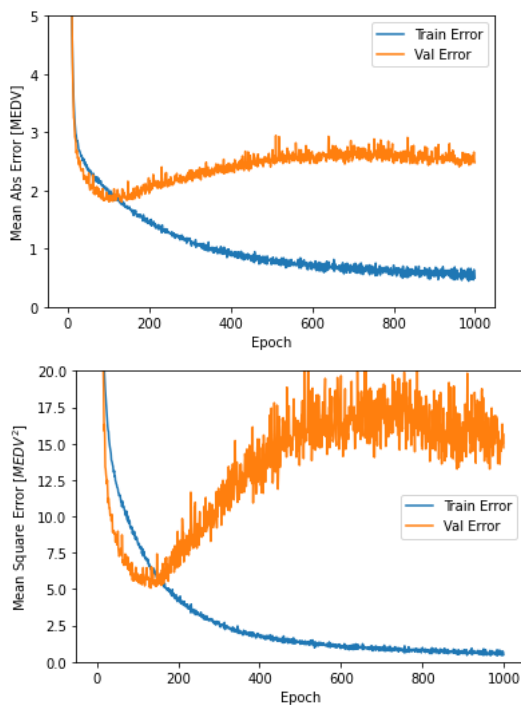
	loss	mae	mse	val_loss	val_mae	val_mse	epoch
<b>995</b>	0.693003	0.567716	0.693003	15.031659	2.488642	15.031659	995
<b>996</b>	0.437564	0.464099	0.437564	16.344011	2.586914	16.344011	996
<b>997</b>	0.727947	0.640294	0.727947	14.969596	2.659240	14.969596	997
<b>998</b>	0.680283	0.568104	0.680283	15.679944	2.519410	15.679944	998
<b>999</b>	0.494136	0.496397	0.494136	14.722670	2.480395	14.722670	999

```
def plot_history(history):
    hist = pd.DataFrame(history.history)
    hist['epoch'] = history.epoch

    plt.figure()
    plt.xlabel('Epoch')
    plt.ylabel('Mean Abs Error [MEDV]')
    plt.plot(hist['epoch'], hist['mae'],
             label='Train Error')
    plt.plot(hist['epoch'], hist['val_mae'],
             label = 'Val Error')
    plt.ylim([0,5])
    plt.legend()

    plt.figure()
    plt.xlabel('Epoch')
    plt.ylabel('Mean Square Error [MEDV^2]')
    plt.plot(hist['epoch'], hist['mse'],
             label='Train Error')
    plt.plot(hist['epoch'], hist['val_mse'],
             label = 'Val Error')
    plt.ylim([0,20])
    plt.legend()
    plt.show()
```

```
plot_history(history)
```





```
model = build_model()

# The patience parameter is the amount of epochs to check for improvement
early_stop = keras.callbacks.EarlyStopping(monitor='val_loss', patience=10)

history = model.fit(normed_train_data, train_labels, epochs=EPOCHS,
                    validation_split = 0.2, verbose=0, callbacks=[early_stop, PrintDot()])
```

```
.....

plot_history(history)
```



```
loss, mae, mse = model.evaluate(normed_test_data, test_labels, verbose=2)
```


```
print("Testing set Mean Abs Error: {:.2f} MEDV".format(mae))
```

```
5/5 - 0s - loss: 10.2339 - mae: 2.3418 - mse: 10.2339 - 30ms/epoch - 6ms/step
Testing set Mean Abs Error: 2.34 MEDV
```

```
test_predictions = model.predict(normed_test_data).flatten()
train_predictions = model.predict(normed_train_data).flatten()
```

```
plt.scatter(test_labels, test_predictions)
plt.xlabel('True Values [MEDV]')
plt.ylabel('Predictions [MEDV]')
plt.axis('equal')
plt.axis('square')
plt.xlim([0,plt.xlim()[1]])
plt.ylim([0,plt.ylim()[1]])
_ = plt.plot([-100, 100], [-100, 100])
```

```
5/5 [=====] - 0s 3ms/step
12/12 [=====] - 0s 2ms/step
```



```
error = test_predictions - test_labels
plt.hist(error, bins = 25)
plt.xlabel("Prediction Error [MEDV]")
_ = plt.ylabel("Count")
```

```
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
mse = mean_squared_error(test_labels, test_predictions)
print('Mean Squared Error: ',mse)
mae = mean_absolute_error(test_labels, test_predictions)
print('Mean Absolute Error: ',mae)
rsq = r2_score(train_labels,train_predictions) #R-Squared on the training data
print('R-square, Training: ',rsq)
rsq = r2_score(test_labels,test_predictions) #R-Squared on the testing data
print('R-square, Testing: ',rsq)
```

```
Mean Squared Error: 11.526060494195113
Mean Absolute Error: 2.3802721861161684
R-square, Training: 0.8939664748792459
R-square, Testing: 0.8835458931882735
```

